

So, You Think You Can Write an EVM Decompiler?

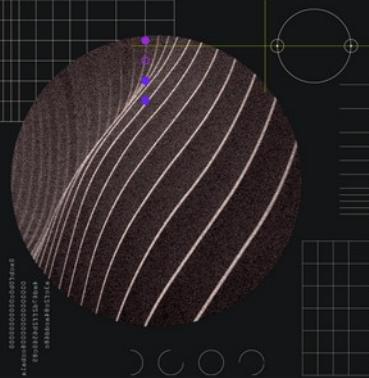
Yannis Smaragdakis (**DEDAUB** , U.Athens)

0xcbfb0b0c000000000000
00000000000000000000
4e469422b15b508ccba1a
a3cf22884dd6981c99d085



Background + Context

- We (Dedaub) have a long-running public decompiler
 - *Gigahorse*
 - app.dedaub.com
 - used by many hundreds every day
 - 10,000 registered users
 - algorithms documented in several research publications
[*ICSE'19, OOPSLA'20, OOPSLA'22, ISSTA'25, ICSE'26*]
 - open source for the most part, source-level proprietary



Point of the Talk

- What is hard about writing an EVM decompiler?
 - how do we do it? Is this the only way?
 - If you already know this stuff:
 - control-flow recovery
 - function recognition
 - memory modeling
 - storage modeling
- Good news: you won't need to understand algorithms!
 - you can do it!
- Bad news: at key points, I need your full attention, so you can understand the problem, from examples

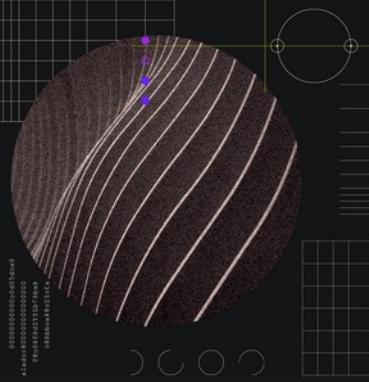




Decompilation

The Ethereum Virtual Machine (EVM)

- Stack based:
 - No concept of variables
- Operations on 256 bit integers
 - No type information
- No abstractions of public functions/entry points or private methods
- JUMP targets can be “dynamic”
- Compilers optimize code in order to reduce code size:
 - Reuse of low-level code blocks when possible



Simple Smart Contract (Solidity)

```
interface IERC20
{ function transfer(address to, uint256 value) external returns (bool); }
contract Contract {
    uint256 defaultFee;

    function simpleTransfer(address tok, address to, uint256 amt) external
    { IERC20(tok).transfer(to, amt); }

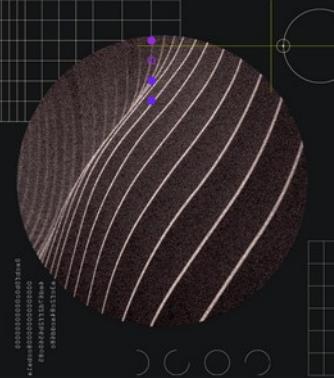
    function transWFee(address tok, address to, uint256 amt,
                      uint256 feeA, uint256 feeB) external
    { IERC20(tok).transfer(to, amt - defaultFee - feeA - feeB); }
    /* makes 3 private function calls! */
}
```



So, You Think You Can Write an EVM Decomplier?

Simple Smart Contract (compiled)

```
608060405234801561000f575f80fd5b5060043610610034575f3560e01c806312e494061461003857806387d7a5f414610
054575b5f80fd5b610052600480360381019061004d9190610227565b610070565b005b61006e6004803603810190610069
9190610277565b6100f1565b005b8273ffffffffffffffffff1663a9059chb83836040518363f
fffffffff1660e01b81526004016100ab92919061030c565b6020604051808303815f875af11580156100c7573d5f803e3d5f
fd5b505050506040513d601f19601f820116820180604052508101906100eb9190610368565b50505050565b8473ffff
ffffffffff1663a9059ccb8583855f548861011d91906103c0565b61012791906103c0565b61
013191906103c0565b6040518363ffff1660e01b815260040161014e92919061030c565b6020604051808303815f875
af115801561016a573d5f803e3d5ffd5b5050506040513d601f19601f8201168201806040525081019061018e91906103
68565b5050505050565b5f80fd5b5f73ffffffffffff82169050919050565b5f6101c
38261019a565b9050919050565b6101d3816101b9565b81146101dd575f80fd5b50565b5f813590506101ee816101ca565b
92915050565b5f819050919050565b610206816101f4565b8114610210575f80fd5b50565b5f81359050610221816101fd5
65b92915050565b5f805f6060848603121561023e5761023d610196565b5b5f61024b868287016101e0565b935050602061
025c868287016101e0565b925050604061026d86828701610213565b9150509250925092565b5f805f805f60a0868803121
56102905761028f610196565b5b5f61029d888289016101e0565b95505060206102ae888289016101e0565b945050604061
02bf88828901610213565b93505060606102d088828901610213565b92505060806102e188828901610213565b915050929
5509295909350565b6102f7816101b9565b82525050565b610306816101f4565b82525050565b5f60408201905061031f5f
8301856102ee565b61032c60208301846102fd565b93925050565b5f8115159050919050565b61034781610333565b811
4610351575f80fd5b50565b5f815190506103628161033e565b92915050565b5f6020828403121561037d5761037c610196
565b5b5f61038a84828501610354565b91505092915050565b7f4e487b71000000000000000000000000000000000000000000000000
00000000000000000005f52601160045260245ffd5b5f6103ca826101f4565b91506103d5836101f4565b9250828203905081
8111156103ed576103ec610393565b5b9291505056fea2646970667358221220e987f5dd559089603fe79b9c0506b8440a1
a70c32d013a6a038c780a82de25d064736f6c63430008190033
```



Simple EVM program (Disassembled): 710 loc

0x0: PUSH1 0x80	0x2b: DUP1	...
0x2: PUSH1 0x40	0x2c: PUSH4	0x131: JUMPDEST
0x4: MSTORE	0x31: EQ	0x132: DUP3
0x5: CALLVALUE	0x32: PUSH2	0x133: MSTORE
0x6: DUP1	0x35: JUMPI	0x134: POP
0x7: ISZERO	0x36: JUMPDEST	0x135: POP
0x8: PUSH2 0x10	0x37: PUSH1	0x136: JUMP
0xb: JUMPI	0x39: DUP1	0x137: JUMPDEST
0xc: PUSH1 0x0	0x3a: REVERT	0x138: PUSH1 0x0
0xe: DUP1	0x3b: JUMPDEST	0x13a: PUSH1 0x20
0xf: REVERT	0x3c: PUSH2	0x13c: DUP3
0x10: JUMPDEST	0x3f: PUSH2	0x13d: ADD
0x11: POP	0x42: JUMP	0x13e: SWAP1
0x12: PUSH1 0x4	0x43: JUMPDEST	0x13f: POP
0x14: CALLDATASIZE	0x44: STOP	0x140: PUSH2 0x14c
0x15: LT	0x45: JUMPDEST	0x143: PUSH1 0x0
0x16: PUSH2 0x36	0x46: PUSH2	0x145: DUP4
0x19: JUMPI	0x49: PUSH2	0x146: ADD
0x1a: PUSH1 0x0	0x4c: JUMP	0x147: DUP5
0x1c: CALLDATALOAD	0x4d: JUMPDEST	0x148: PUSH2 0x128
0x1d: PUSH1 0xe0	0x4e: PUSH1	0x14b: JUMP
0x1f: SHR	0x50: MLOAD	0x14c: JUMPDEST
0x20: DUP1	0x51: PUSH2	0x14d: SWAP3
0x21: PUSH4 0x12e49406	0x54: SWAP2	0x14e: SWAP2
0x26: EQ	0x55: SWAP1	0x14f: POP
0x27: PUSH2 0x3b	0x56: PUSH2	0x150: POP
0x2a: JUMPI	0x59: JUMP	0x151: JUMP



Simple EVM program (Disassembled): 710 loc

contract
entry-point

```
0x0: PUSH1 0x80
0x2: PUSH1 0x40
0x4: MSTORE
0x5: CALLVALUE
0x6: DUP1
0x7: ISZERO
0x8: PUSH2 0x10
0xb: JUMPI
0xc: PUSH1 0x0
0xe: DUP1
0xf: REVERT
0x10: JUMPDEST
0x11: POP
0x12: PUSH1 0x4
0x14: CALLDATASIZE
0x15: LT
0x16: PUSH2 0x36
0x19: JUMPI
0x1a: PUSH1 0x0
0x1c: CALLDATALOAD
0x1d: PUSH1 0xe0
0x1f: SHR
0x20: DUP1
0x21: PUSH4 0x12e49406
0x26: EQ
0x27: PUSH2 0x3b
0x2a: JUMPI
```

```
0x2b: DUP1
0x2c: PUSH4
0x31: EQ
0x32: PUSH2
0x35: JUMPI
0x36: JUMPDEST
0x37: PUSH1
0x39: DUP1
0x3a: REVERT
0x3b: JUMPDEST
0x3c: PUSH2
0x3f: PUSH2
0x42: JUMP
0x43: JUMPDEST
0x44: STOP
0x45: JUMPDEST
0x46: PUSH2
0x49: PUSH2
0x4c: JUMP
0x4d: JUMPDEST
0x4e: PUSH1
0x50: MLOAD
0x51: PUSH2
0x54: SWAP2
0x55: SWAP1
0x56: PUSH2
0x59: JUMP
```

```
... 0x131: JUMPDEST
0x132: DUP3
0x133: MSTORE
0x134: POP
0x135: POP
0x136: JUMP
0x137: JUMPDEST
0x138: PUSH1 0x0
0x13a: PUSH1 0x20
0x13c: DUP3
0x13d: ADD
0x13e: SWAP1
0x13f: POP
0x140: PUSH2 0x14c
0x143: PUSH1 0x0
0x145: DUP4
0x146: ADD
0x147: DUP5
0x148: PUSH2 0x128
0x14b: JUMP
0x14c: JUMPDEST
0x14d: SWAP3
0x14e: SWAP2
0x14f: POP
0x150: POP
0x151: JUMP
```

0x1a: PUSH1 0x0
0x1c: CALLDATALOAD
0x1d: PUSH1 0xe0
0x1f: SHR
0x20: DUP1
0x21: PUSH4 0x12e49406
0x26: EQ
0x27: PUSH2 0x3b
0x2a: JUMPI

0x45: JUMPDEST
0x46: PUSH2 0x4d
0x49: PUSH2 0x104
0x4c: JUMP

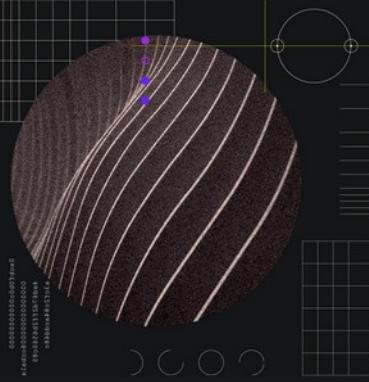
0x14d: SWAP3
0x14e: SWAP2
0x14f: POP
0x150: POP

0x151: JUMP

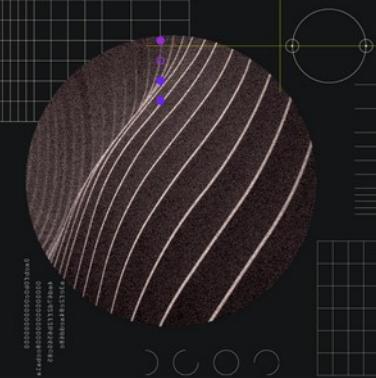


Decompilation Definition

The recovery of a structured intermediate representation (three-address code) from very low-level bytecode.

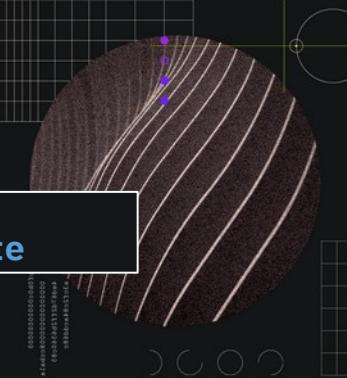


Gigahorse Three-Address-Code (TAC)



So, You Think You Can Write an EVM Decompiler?

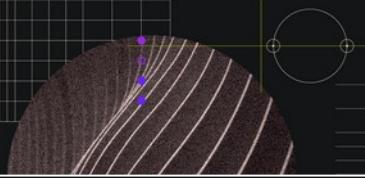
Gigahorse Three-Address-Code (TAC)



Begin block 0xc7

So, You Think You Can Write an EVM Decompiler?

Gigahorse Three-Address-Code (TAC)



```
function @simpleTransfer_59(v70arg0, v70arg1, v70arg2, v70arg3)
```

```
=====
0xac: vac(0x20) = CONST
0xae: vae(0x40) = CONST
0xb0: vb0 = MLOAD vae(0x40)
0xb3: vb3 = SUB vaa_0, vb0
0xb5: vb5(0x0) = CONST
0xb7: vb7 = GAS
0xb8: vb8 = CALL vb7, v87, vb5(0x0), vb0,
0xb9: vb9 = ISZERO vb8
0xbb: vbb = TSZERO vba
```

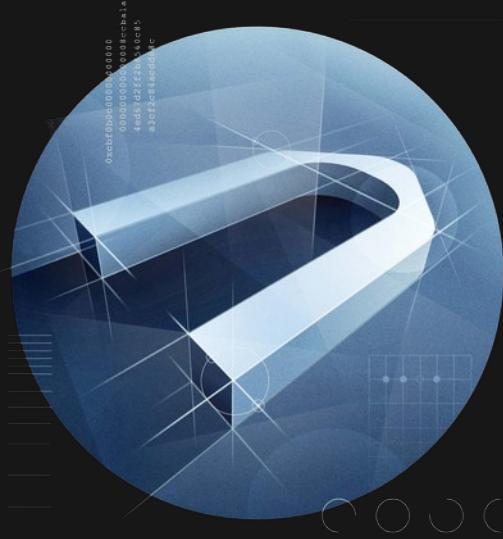
0xbf: JUMPI vbc(0xc7), vbb

```
Begin block 0xc0
prev=[0xab], succ=[]
=====
0xc0: vc0 = RETURNDATASIZE
0xc1: vc1(0x0) = CONST
0xc3: RETURNDATACOPY vc1(0x0), vc1(0x0), vc0
0xc4: vc4 = RETURNDATASIZE
```

0xaaa: vaa_0 = CALLPRIVATE va7(0x30c), va0, v70arg0, v70arg1, va1(0xab)

Begin block 0xc7





Difficulties #1, #2: Control Flow + Functions



So, You Think You Can Write an EVM Decompiler?

Simple EVM program (Disassembled): 710 loc

contract
entry-point

```
0x0: PUSH1 0x80
0x2: PUSH1 0x40
0x4: MSTORE
0x5: CALLVALUE
0x6: DUP1
0x7: ISZERO
0x8: PUSH2 0x10
0xb: JUMPI
0xc: PUSH1 0x0
0xe: DUP1
0xf: REVERT
0x10: JUMPDEST
0x11: POP
0x12: PUSH1 0x4
0x14: CALLDATASIZE
0x15: LT
0x16: PUSH2 0x36
0x19: JUMPI
0x1a: PUSH1 0x0
0x1c: CALLDATALOAD
0x1d: PUSH1 0xe0
0x1f: SHR
0x20: DUP1
0x21: PUSH4 0x12e49406
0x26: EQ
0x27: PUSH2 0x3b
0x2a: JUMPI
```

```
0x2b: DUP1
0x2c: PUSH4
0x31: EQ
0x32: PUSH2
0x35: JUMPI
0x36: JUMPDEST
0x37: PUSH1
0x39: DUP1
0x3a: REVERT
0x3b: JUMPDEST
0x3c: PUSH2
0x3f: PUSH2
0x42: JUMP
0x43: JUMPDEST
0x44: STOP
0x45: JUMPDEST
0x46: PUSH2
0x49: PUSH2
0x4c: JUMP
0x4d: JUMPDEST
0x4e: PUSH1
0x50: MLOAD
0x51: PUSH2
0x54: SWAP2
0x55: SWAP1
0x56: PUSH2
0x59: JUMP
...
```

```
... 0x131: JUMPDEST
0x132: DUP3
0x133: MSTORE
0x134: POP
0x135: POP
0x136: JUMP
0x137: JUMPDEST
0x138: PUSH1 0x0
0x13a: PUSH1 0x20
0x13c: DUP3
0x13d: ADD
0x13e: SWAP1
0x13f: POP
0x140: PUSH2 0x14c
0x143: PUSH1 0x0
0x145: DUP4
0x146: ADD
0x147: DUP5
0x148: PUSH2 0x128
0x14b: JUMP
0x14c: JUMPDEST
0x14d: SWAP3
0x14e: SWAP2
0x14f: POP
0x150: POP
0x151: JUMP
...
```

0x1a: PUSH1 0x0
0x1c: CALLDATALOAD
0x1d: PUSH1 0xe0
0x1f: SHR
0x20: DUP1
0x21: PUSH4 0x12e49406
0x26: EQ
0x27: PUSH2 0x3b
0x2a: JUMPI

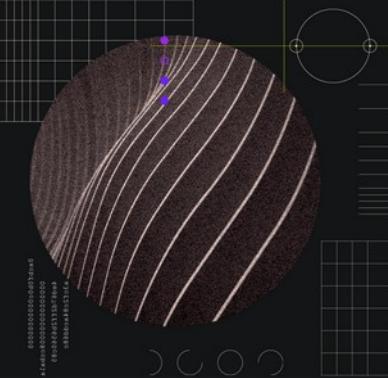
0x45: JUMPDEST
0x46: PUSH2 0x4d
0x49: PUSH2 0x104
0x4c: JUMP

0x14d: SWAP3
0x14e: SWAP2
0x14f: POP
0x150: POP

0x151: JUMP

Heart of the Problem

- **Execution of the code determines structure of the code!**
- Need work to tell if a **JUMP** is:
 - **if**
 - **for**
 - function call
 - return
- Secondarily: *does the top of the stack at this instruction contain a jump label, or a run-time value?*



0x45: JUMPDEST
0x46: PUSH2 **0x4d**
0x49: PUSH2 **0x104**
0x4c: JUMP

0x14d: SWAP3

0x14e: SWAP2

0x14f: POP

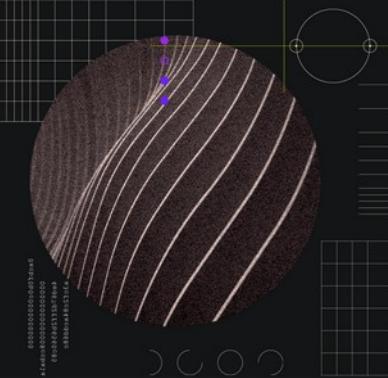
0x150: POP

0x151: JUMP



"Obvious" solution

- Execution of the code determines structure of the code?
- **Just execute the code to discover its structure!**
 - several decompilers have tried to do just that: *symbolic execution*
 - Porosity
 - Panoramix
 - Heimdall-rs
 - *that's not what we do*



```
0x45: JUMPDEST  
0x46: PUSH2 0x4d  
0x49: PUSH2 0x104  
0x4c: JUMP
```

```
0x14d: SWAP3
```

```
0x14e: SWAP2
```

```
0x14f: POP
```

```
0x150: POP
```

```
0x151: JUMP
```

Comparison to Heimdall-rs: Completeness

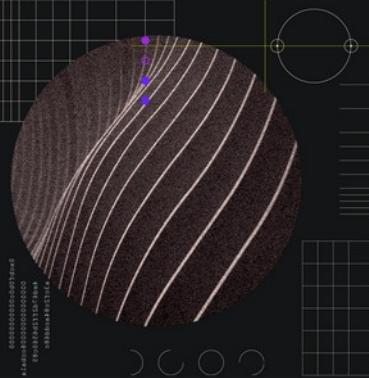
Yul Dataset

	Unique External Calls		Unique Events	
Gigahorse	13600	+	13661	+
Heimdall-rs	9841	-	9505	-



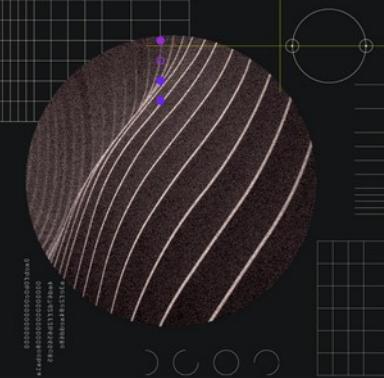
Problem with "Just Execute the Code"

- Code has exponential (or infinite) number of paths
 - how can you cover all?
 - only need to simulate the parts of the code that involve jump labels, not all values!
- Compiler has performed very complex optimizations, merged code from far-away source positions



Problem with "Just Execute the Code"

- Code has exponential (or infinite) number of paths
 - how can you cover all?
 - only need to simulate the parts of the code that involve jump labels, not all values!
- Compiler has performed very complex optimizations, merged code from far-away source positions
 - *is that a label at the top of the stack?*
 - *depends!*



0x14c: JUMPDEST

0x14d: SWAP3

0x14e: SWAP2

0x14f: POP

0x150: POP

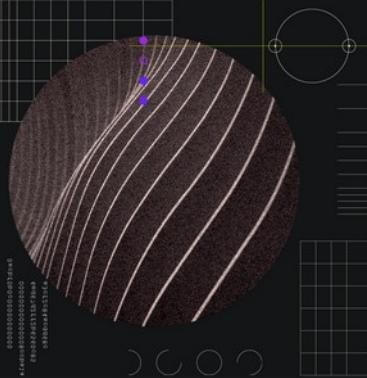
0x151: JUMP



So, You Think You Can Write an EVM Decompiler?

What We Do Instead: *Static Analysis*

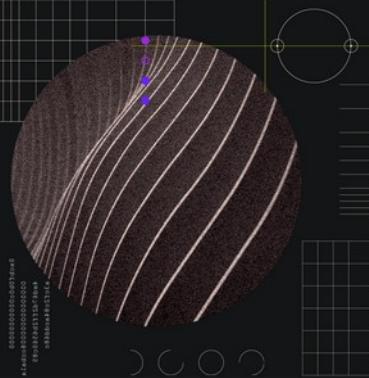
- Approximation of what the code does for all possible executions
 - how to distinguish ones that do things differently?
 - "context sensitivity"
 - classify executions into a finite (but large) number of bins
 - use the same algorithm to also tell us which parts of the code involve jump labels positions



Our Algorithm (*slightly simplified*)

When calling a function do we need to remember where it was called from?

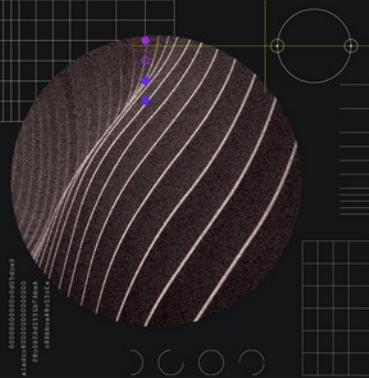
Yes.



Our Algorithm (*slightly simplified*)

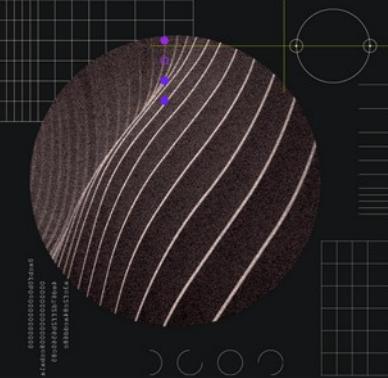
When a function returns do we need to remember it was ever called?

No.



Similar Complications: Recognizing (Private) Functions

- How many arguments are passed in this call? Of what types? How many values are returned?
 - maybe it depends?
 - some call site may be ignoring some returns?
 - EVM has no well-structured stack
- Need again complex algorithms



0x45:	JUMPDEST
0x46:	PUSH2
0x49:	PUSH2
0x4c:	JUMP
	0x4d
	0x104

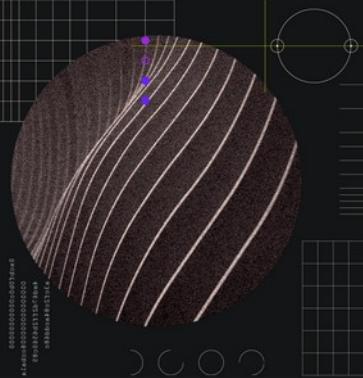
0x0000000000000000
00000000000000000000000000000000
4-4d7-02f22b510c45
a1cf2c44cdde4c



Difficulties #3, #4: Memory Modeling, Storage Modeling

EVM Data Stores, through Solidity

- Two kinds of stores:
 - **Storage:** Persistent, kept on the blockchain's state
 - **Memory:** Volatile, per transaction
- Memory is used by the compiler as a scratchpad for all computations of data with undetermined length
 - arrays, strings



Solidity Smart Contract Example

```
contract Example{  
    string onStorage;  
  
    function setIt(string memory newStr) public {  
        onStorage = newStr;  
    }  
  
    function getHash() public view returns (bytes32) {  
        return keccak256(onStorage);  
    }  
}
```

On persistent storage

Argument passed through memory

Return passed through memory

Keccak hashes memory contents



“Memory” hides a lot of implicit computation!

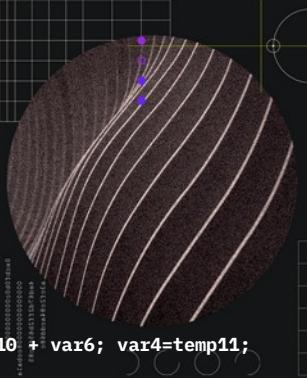
```
mapping(string => string) mTokens; ...
function getToken(string pDocumentHash) returns(string)
{ return mTokens[pDocumentHash]; }
```

becomes:

```
function getToken(var arg0) returns (var r0) {
    var var0 = 0x053b; var0 = func_06C6(); var var1 = 0x02;
    var temp0 = arg0; var var2 = temp0;
    var var3 = memory[0x40:0x60]; var var4 = var3;
    var var5 = var2 + 0x20; var var6 = memory[var2:var2 + 0x20];
    var var7 = var6; var var8 = var4; var var9 = var5;
    if (var7 < 0x20) {
        label_0573:
        var temp1 = 0x0100 ** (0x20 - var7) - 0x01; var temp2 = var8;
        memory[temp2:temp2 + 0x20] = (memory[var9:var9 + 0x20] &
        ~temp1) | (memory[temp2:temp2 + 0x20] & temp1);
        var temp3 = var6 + var4;
        memory[temp3:temp3 + 0x20] = var1;
        var temp4 = memory[0x40:0x60];
        var temp5 = keccak256(memory[temp4:temp4+(temp3+0x20)-temp4]);
        var temp6 = storage[temp5];
        var temp7 = (!(temp6 & 0x01) * 0x0100 - 0x01 & temp6) / 0x02;
        var temp8 = memory[0x40:0x60];
        memory[0x40:0x60] = temp8 + (temp7+0x1f) / 0x20 * 0x20 + 0x20;
        var1 = temp8; var2 = temp5; var3 = temp7;
        memory[var1:var1 + 0x20] = var3;
        var4 = var1 + 0x20; var5 = var2;
        var temp9 = storage[var5];
        var6 = (!(temp9 & 0x01) * 0x0100 - 0x01 & temp9) / 0x02;
    }
    if (!var6) {
        label_063A:
        return var1;
    } else if (0x1f < var6) {
        var temp10=var4; var temp11 = temp10 + var6; var4=temp11;
        memory[0x00:0x20] = var5;
        var temp12 = keccak256(memory[0x00:0x20]);
        memory[temp10:temp10 + 0x20] = storage[temp12];
        var5 = temp12 + 0x01; var6 = temp10 + 0x20;
        if (var4 <= var6) { goto label_0631; }
    }
    label_061D:
    var temp13 = var5; var temp14 = var6;
    memory[temp14:temp14 + 0x20] = storage[temp13];
    var5 = temp13 + 0x01; var6 = temp14 + 0x20;
    if (var4 > var6) { goto label_061D; }
}
label_0631:
var temp15 = var4; var temp16 = temp15+(var6 - temp15&0x1f);
var6 = temp15; var4 = temp16;
goto label_063A;
} else {
    var temp17 = var4;
    memory[temp17:temp17+0x20] = storage[var5]/0x0100 * 0x0100;
    var4 = temp17 + 0x20; var6 = var6;
    goto label_063A;
}
} else {
    var temp18 = var9; var temp19 = var8;
    memory[temp19:temp19 + 0x20] = memory[temp18:temp18 + 0x20];
    var8 = temp19 + 0x20; var9 = temp18 + 0x20; var7 = var7-0x20;
    if (var7 < 0x20) { goto label_0573; }
    else { goto label_0559; }
}
```

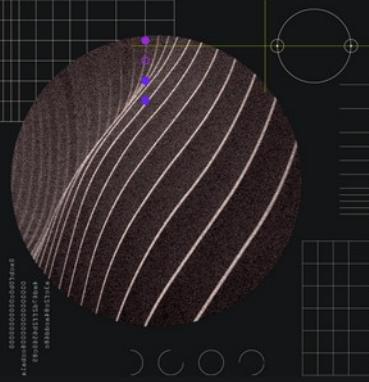


So, You Think You Can Write an EVM Decompiler?



Memory Modeling: A Pretty Hard Problem

- Need to truly reverse engineer fairly arbitrary pointer computation
 - Solidity *free memory pointer*
 - and getting worse





Status of Decompilation

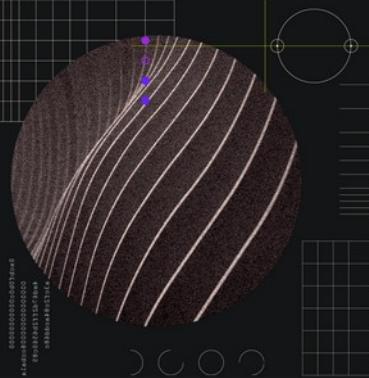


So, You Think You Can Write an EVM Decomplier?

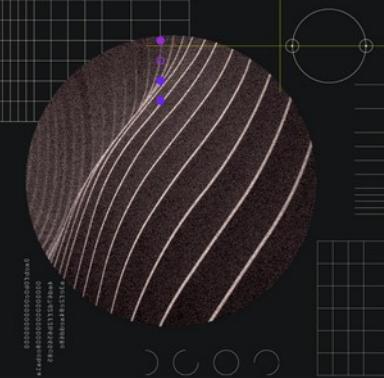
Parting Thoughts: Are the Tough Problems Solved?

Disclaimer: Subjective

- *Difficulty #1:* Control Flow
 - **99%**, you may want to do it differently
- *Difficulty #2:* Function Recovery
 - **95%**
- *Difficulty #3:* Memory Modeling
 - **90%**, may never be 100%
- *Difficulty #4:* Storage Modeling
 - **99+%**: you probably want to imitate what we do [ICSE'26]



Dedaub Gigahorse 3.0



Open source, available at: <https://github.com/nevillegreh/gigahorse-to-olchain>

Deployed on all deployed contracts on many blockchains at:
<https://app.dedaub.com/>

