

18 NOV 2025



# SSA-CFG Yul

A new viaIR backend

clonker / Moritz Hoffmann

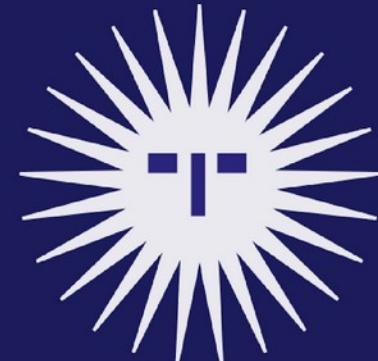




## Performance of current viaIR How to go forward

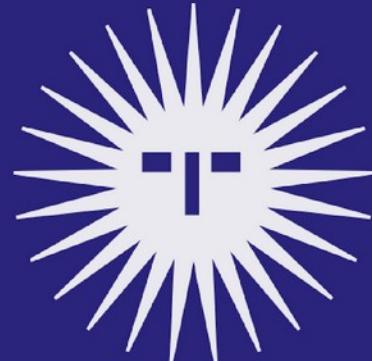
What is SSA-CFG?

Codegen on SSA-CFG Yul  
Liveness and counts  
Stack layout generation  
Stack shuffling  
Terminating paths





# Performance of current viaIR

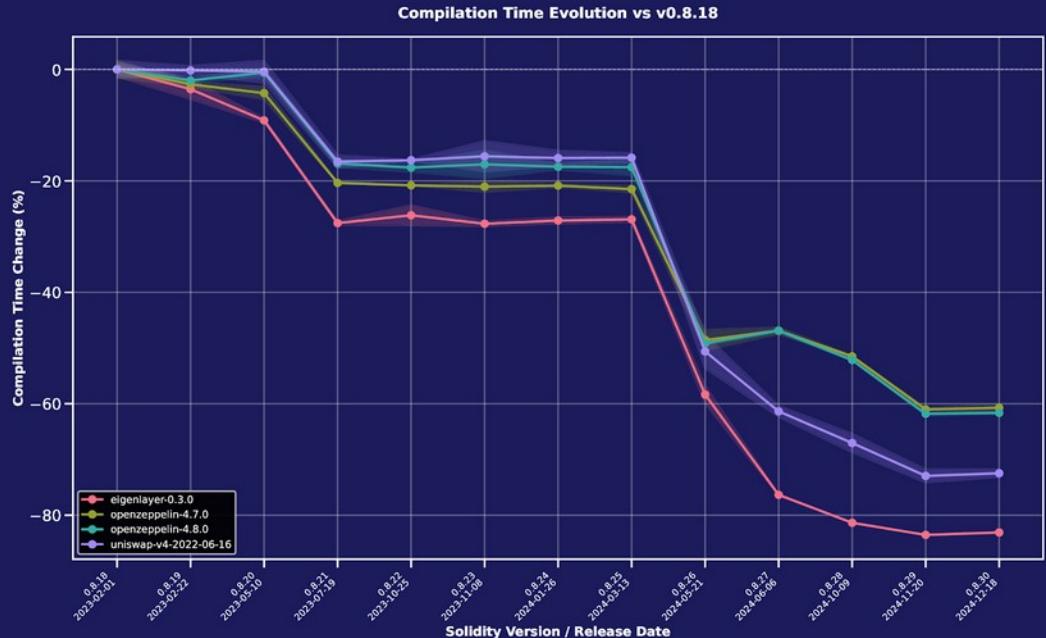




## vialR over the versions and years

on v0.8.30:

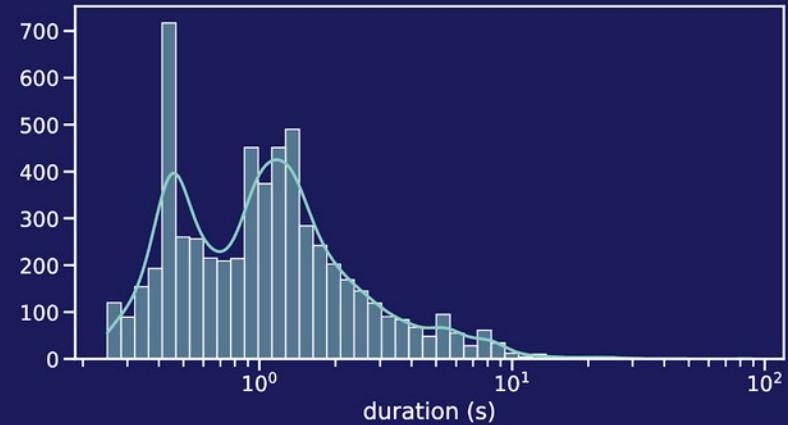
- eigenlayer: ~ 4.4m
- OZ v4.7.0: ~ 24.7s
- OZ v4.8.0: ~ 27.0s





## Gathering some data

- compiled ~6k contracts with latest solc release (v0.8.30)
- 5.5% took longer than 5s
- individual contracts take minutes



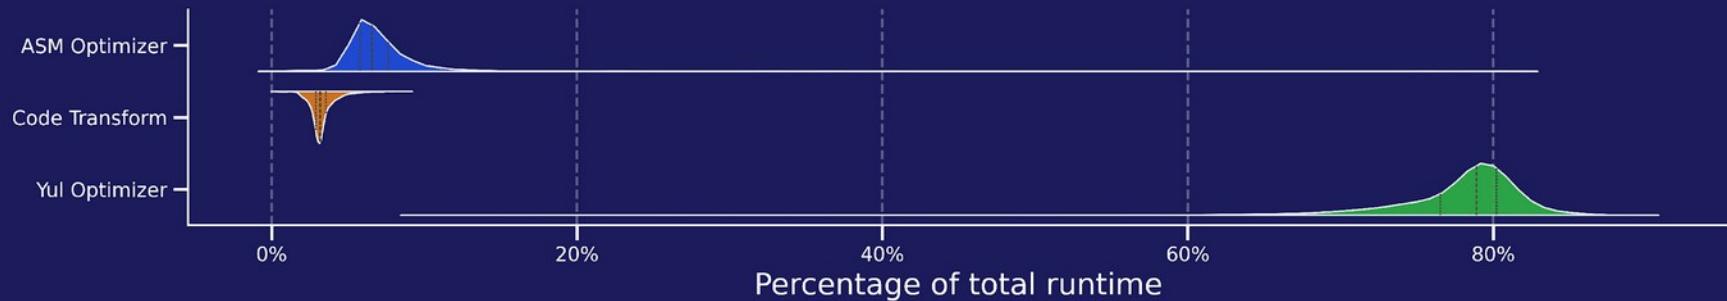
```
select * from compiled_contracts where compiler == 'solc' and
→ version like "0.8.27%"
```

<https://docs.sourcify.dev/docs/repository/sourcify-database/>



## vialR pipeline performance

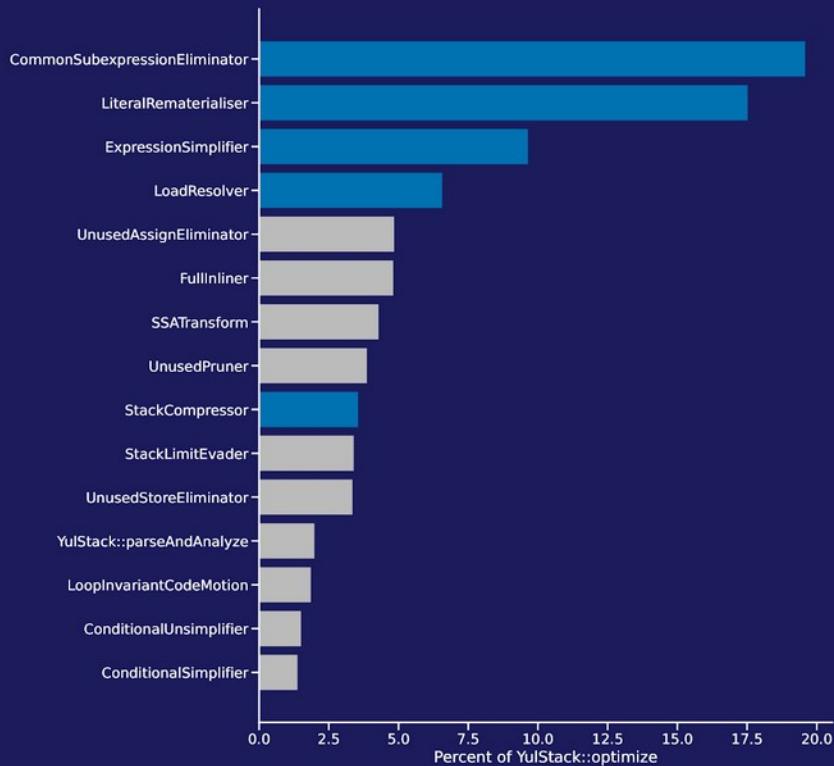
```
solc --via-ir --optimize *args  
CompilerStack::compile  
YulStack::parseAndAnalyze  
YulStack::optimize  
YulStack::assembleEVMWithDeployed  
OptimizedEVMCodeTransform::run  
Assembly::optimize  
bytecode
```





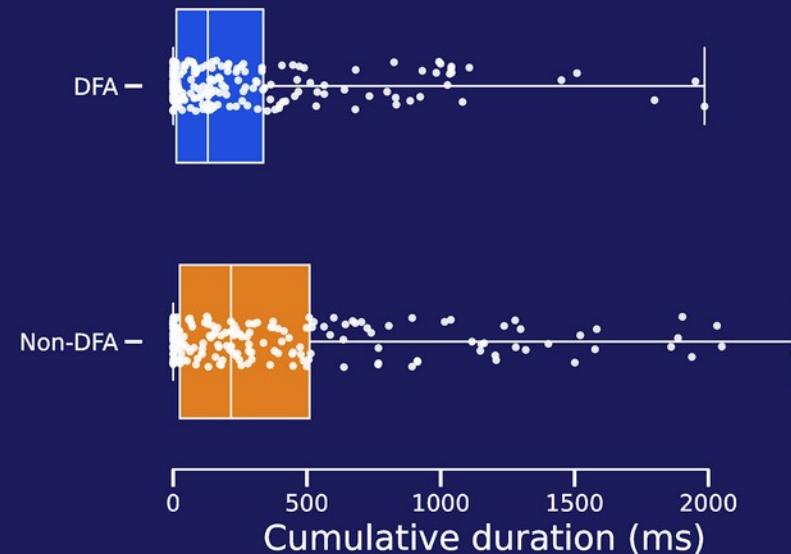
## Yul optimizer performance: Steps depending on data flow

- Strongly depends on structure of input
- For eigenlayer v0.3.0--on my machine 5min--up to 60% of YulStack::optimize in data-flow dependent optimizer passes



## Yul optimizer performance: Steps depending on data flow

- Strongly depends on structure of input
- For eigenlayer v0.3.0--on my machine 5min--up to 60% of YulStack::optimize in data-flow dependent optimizer steps
- But also: Cumulatively 30% more time spent in non data-flow dependent steps



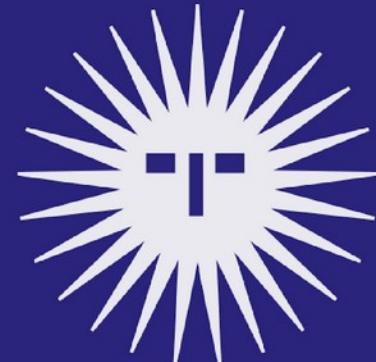


## What we learn from the performance data

- Eigenlayer v0.3.0: takes 5 minutes to compile.
- Bottleneck depends on specific AST, different steps of the optimizer can dominate runtime → need improvements across various steps.
- SSA form enables  $\mathcal{O}(1)$  def-use queries and more efficient data-flow analyses; proven effective in LLVM/GCC.
- Gives opportunity to design and benchmark data structures and algorithms to finally enable fast viaIR compilation.
- Enables easier to maintain codegen and better stack-too-deep handling, too!



# What is SSA-CFG?





## SSA form

- Each variable is assigned exactly once.

### non SSA

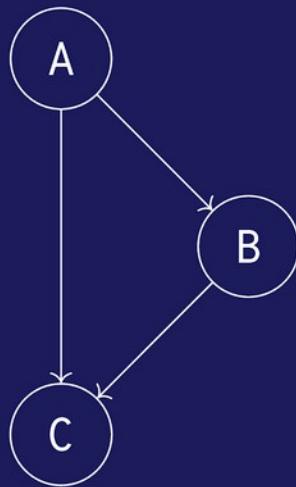
```
let y := 1  
y := 2  
let x := y
```

### SSA

```
let v_1 := 1  
let v_2 := 2  
let v_3 := v_2
```

- Used in many SOTA compilers (LLVM, GCC, glslang, ...)
- Makes data-flow based optimizations more efficient

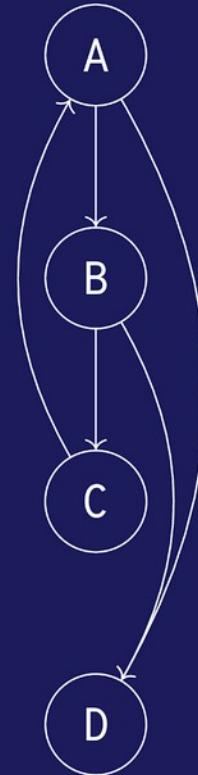
## CFG



if (c) { ; } ;



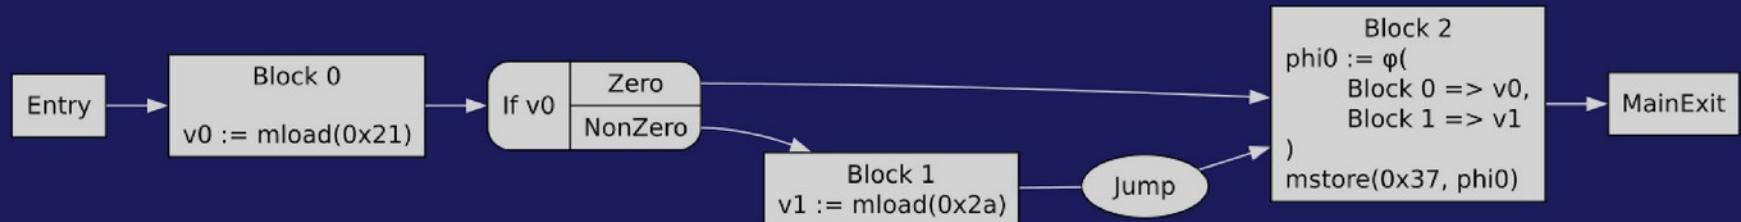
for (;;) { ; } ;





## The Yul SSA-CFG: If

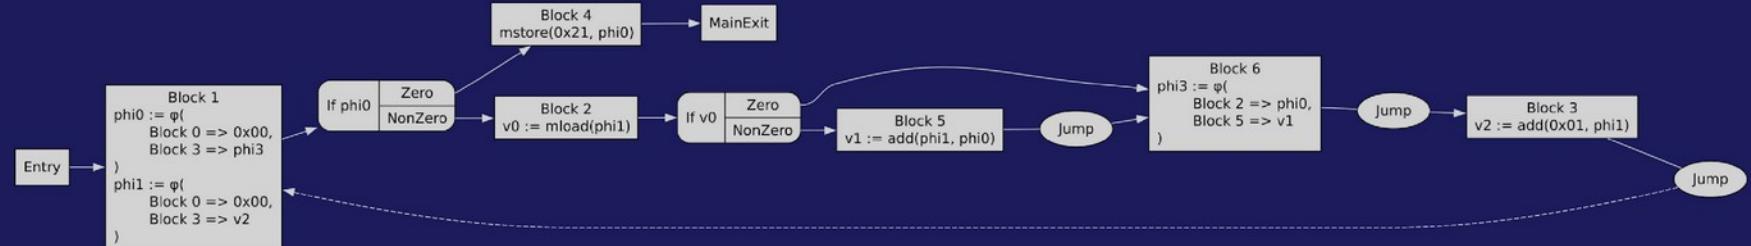
```
{  
    let x := mload(33)  
    if x {  
        x := mload(42)  
    }  
    mstore(x, 55)  
}
```





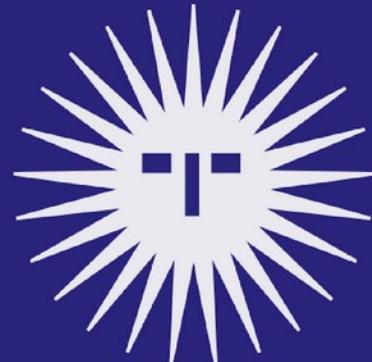
## The Yul SSA-CFG: For

```
{  
    let x  
    for {let i := 0} x {i := add(i, 1)} {  
        if mload(i) {  
            x := add(x, i)  
        }  
    }  
    mstore(x, 33)  
}
```



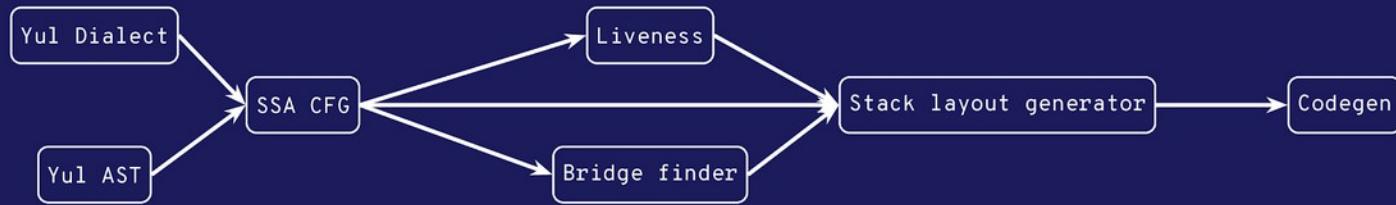


# Codegen on SSA-CFG Yul





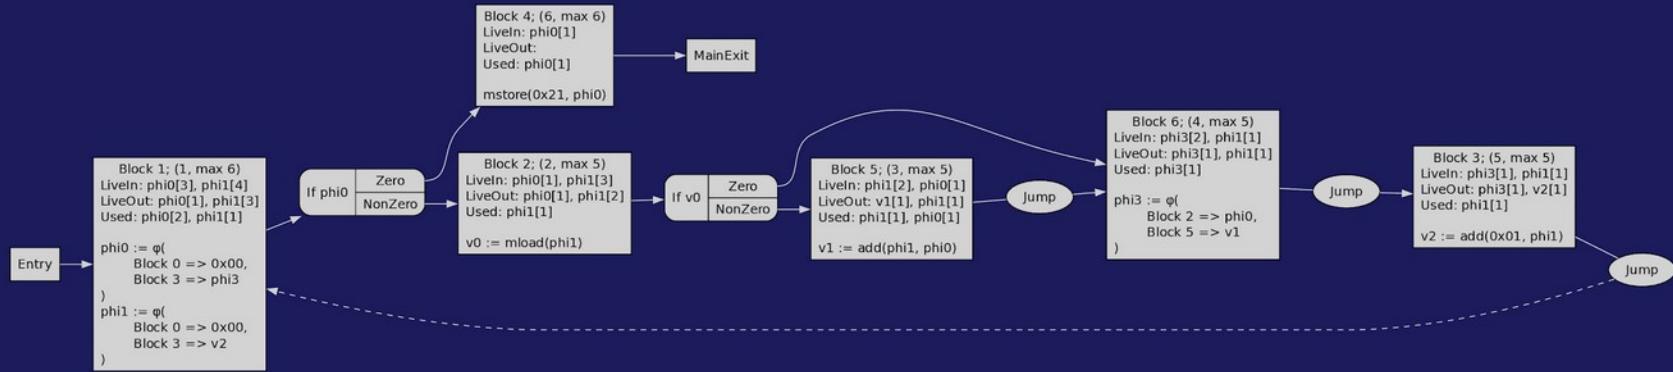
## From Yul to SSA-CFG to bytecode



- Constructing the SSA-CFG based on  
Braun, M. et al. Simple and efficient construction of static single assignment form. In Compiler Construction: 22nd International Conference, CC 2013, ETAPS 2013, Rome, Italy, March 16–24, 2013. Proceedings 22 (pp. 102–122). Springer Berlin Heidelberg.
- Liveness analysis based on  
Rastello, F., & Tichadou, F. B. (Eds.). (2022). SSA-based Compiler Design. Springer.



## Liveness example



Info on

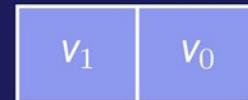
- ... what is live at which point in time,
- ... what is used in a block,
- ... how often something is used downstream.



## Stack layout generation

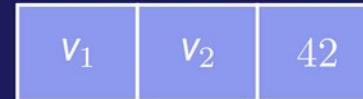
Inside a block, for each operation, assign input layouts.

stackIn



**function**  $f(v\_1, v\_0) \rightarrow v\_4$

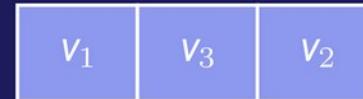
PUSH1 42



**mload**

$v\_2 := mload(v\_0)$

SWAP1



**mload**

$v\_3 := mload(42)$

SWAP1 POP



**add**

$v\_4 := add(v2, v3)$

}



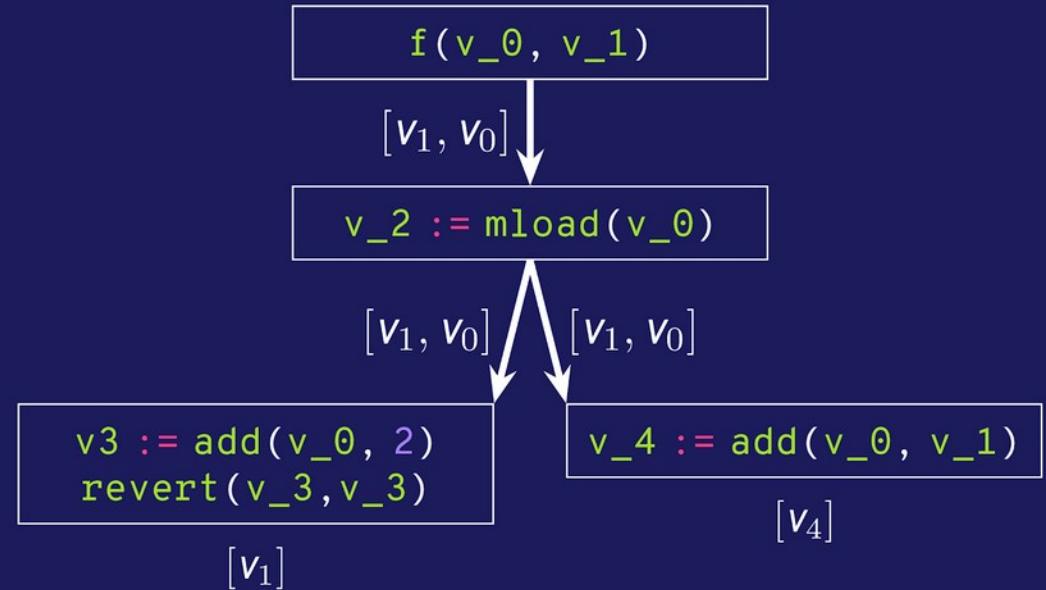
## Stack layout generation

For each block, assign input/output stack layouts.

---

```
function f(a, b) -> c {
    if mload(a) {
        a := add(a, 2)
        revert(a, a)
    }
    c := add(a, b)
}
```

---

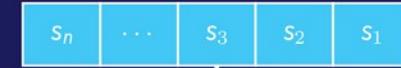




## The greedy operation forward shuffler

- For each operation, prepare an admissible stack
- Each step in the shuffler has to strictly improve the situation

Input S:



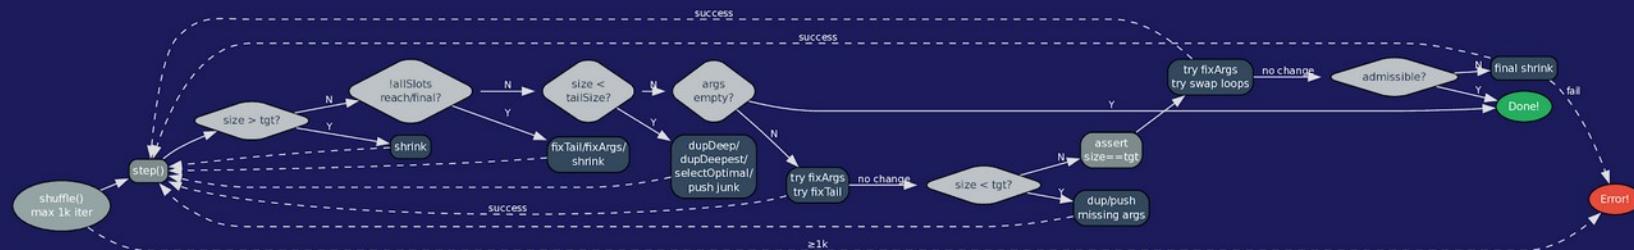
operation  
forward  
shuffler  
(args, liveOut, size)

Target T:



tail(unordered, liveOut)

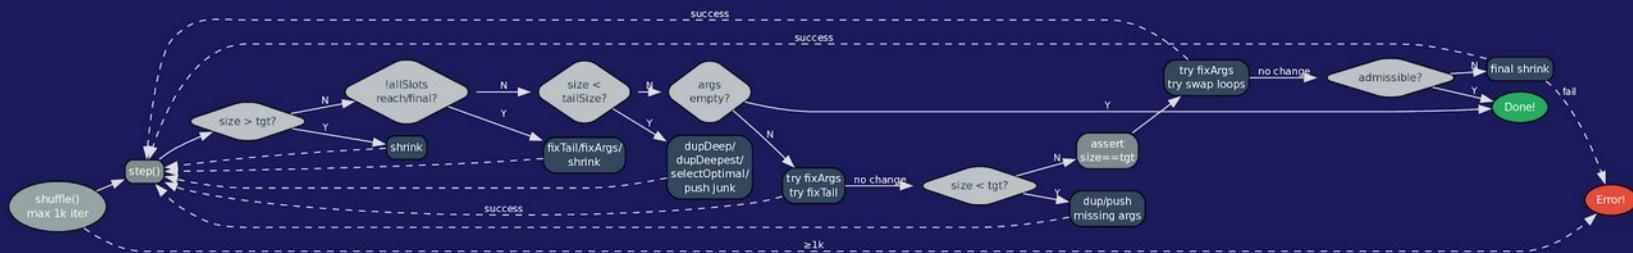
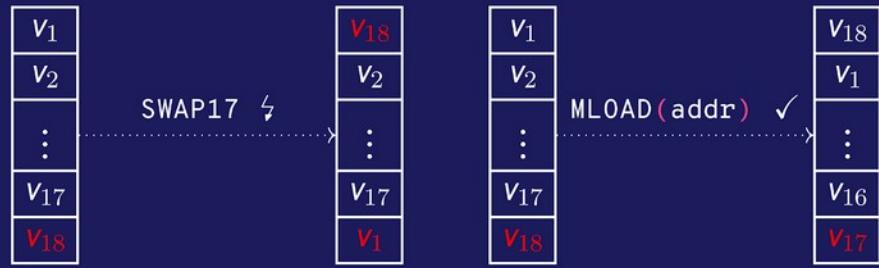
args(ordered)





## The greedy operation forward shuffler

- For each operation, prepare an admissible stack
- Fix stack-too-deep for unreachable operation arguments

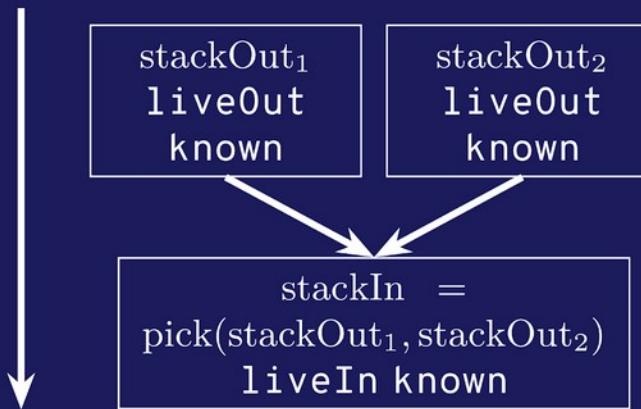




## Comparison to current viaIR stack layout generation

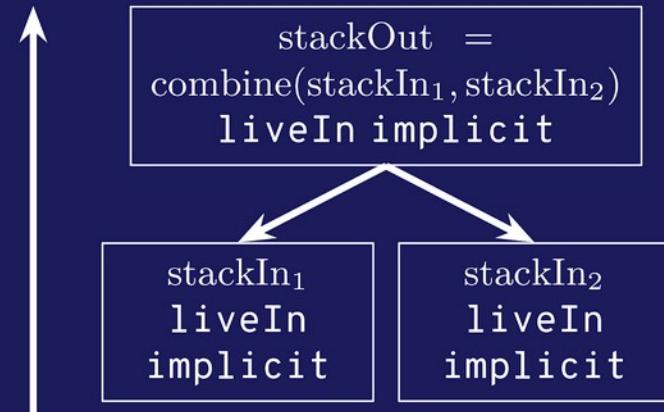
### SSA-CFG viaIR

- top-down stack layout gen
- explicit liveness
- constructive control-flow joins



### current viaIR

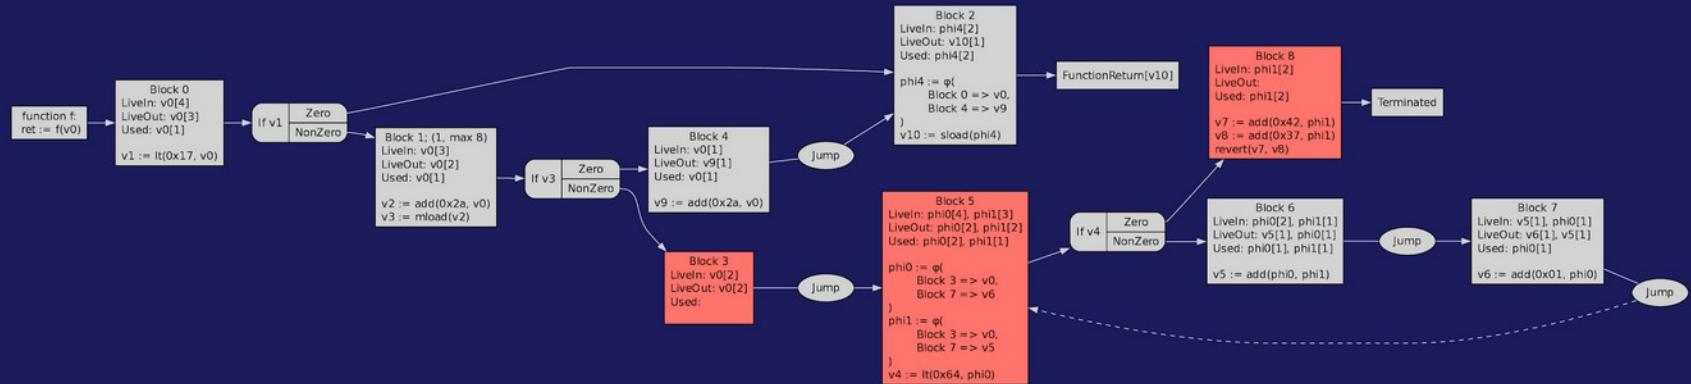
- bottom-up stack layout gen
- implicit liveness
- heap algorithm for control-flow joins





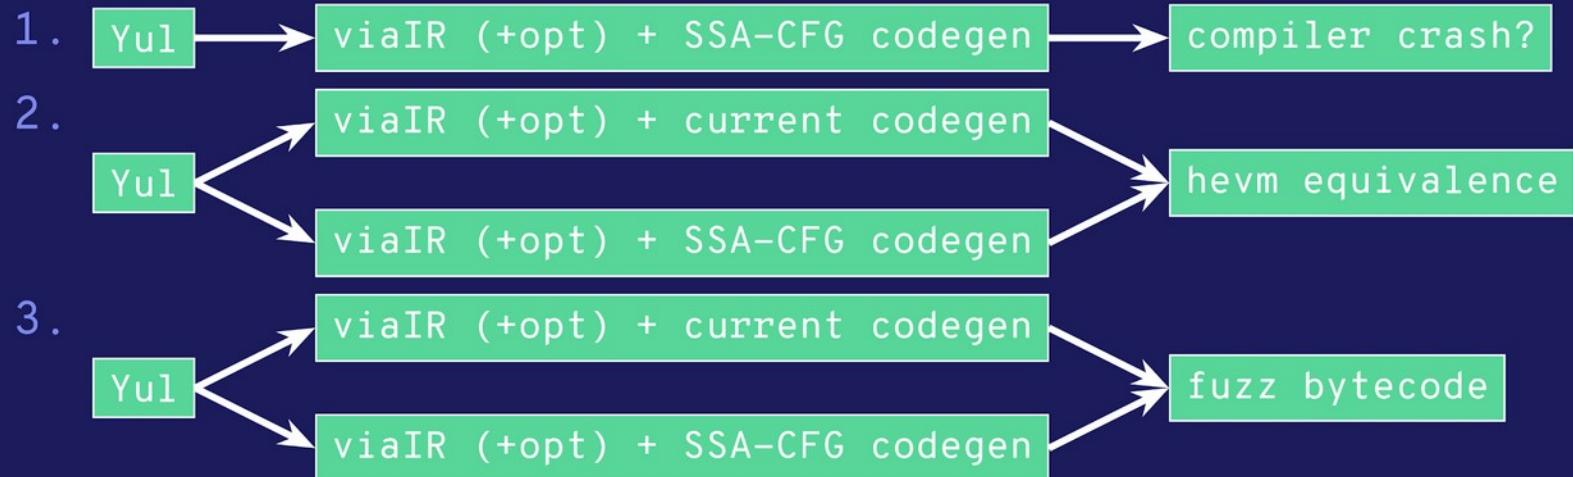
## Terminating paths

- Opposing to user function returns, EVM termination cleans up whatever is left on stack. Saves POPs.
- When do we want to retain superfluous things on stack?
- Explicit in top-down SSA-CFG viaIR!





## (Fuzz-)testing the implementation



Yul may be generated by

- producing random Solidity code,
- producing random Yul code directly,
- or by drawing it from, e.g., Sourcify.



## Outlook

- Initial prototypical implementation without stack-too-deep handling almost ready to be released under experimental flag in solc.
- Bytecode performance on par and often even better than current viaIR.
- Time needed to compile comparable as the heavy optimizer passes are still part of the non-SSA bit of the pipeline.
- Next up:
  - Stack to memory moving
  - Stabilizing the current implementation through (fuzz) testing and extensive reviews
  - Implement data-flow heavy optimizer steps in the SSA-CFG representation to speed up compile time.



# Q&A